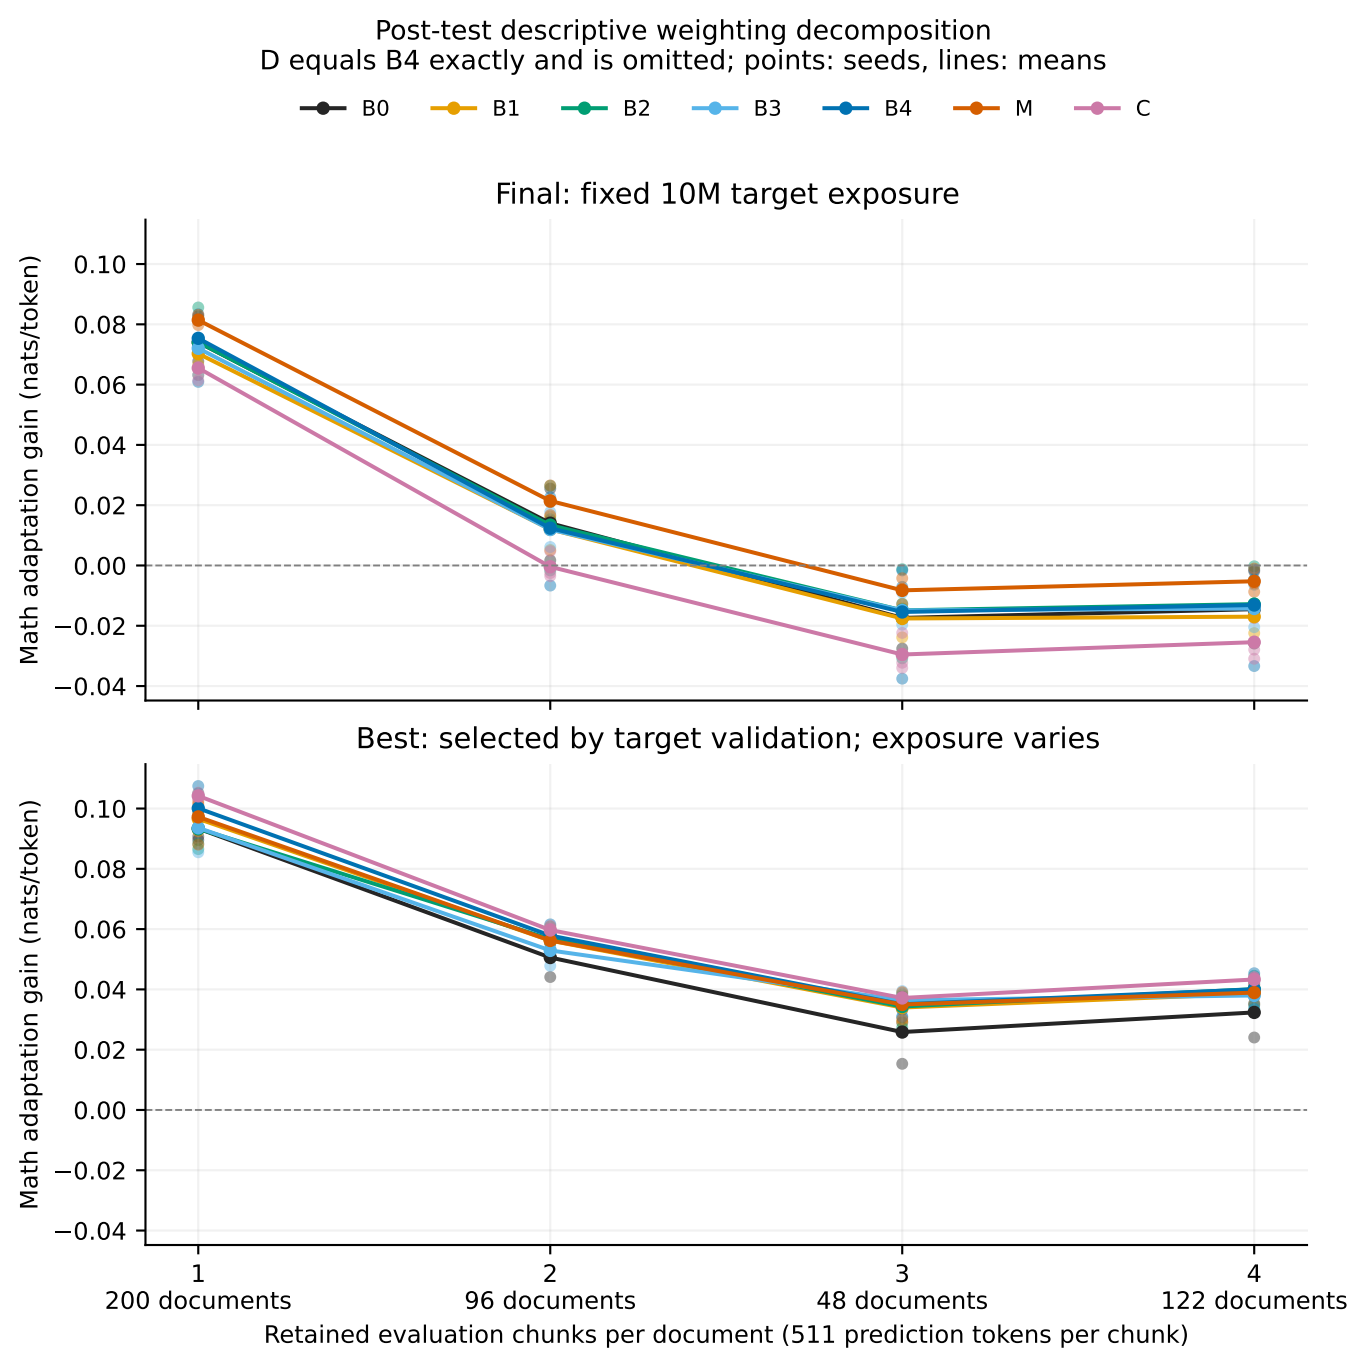
Post-test descriptive weighting decomposition
D equals B4 exactly and is omitted; points: seeds, lines: means
B0
B1
B2
B3
B4
M
C
Final: fixed 10M target exposure
Math adaptation gain (nats/token)
0.10
0.08
0.06
0.04
0.02
0.00
−0.02
−0.04
Best: selected by target validation; exposure varies
Math adaptation gain (nats/token)
0.10
0.08
0.06
0.04
0.02
0.00
−0.02
−0.04
1
200 documents
2
96 documents
3
48 documents
4
122 documents
Retained evaluation chunks per document (511 prediction tokens per chunk)