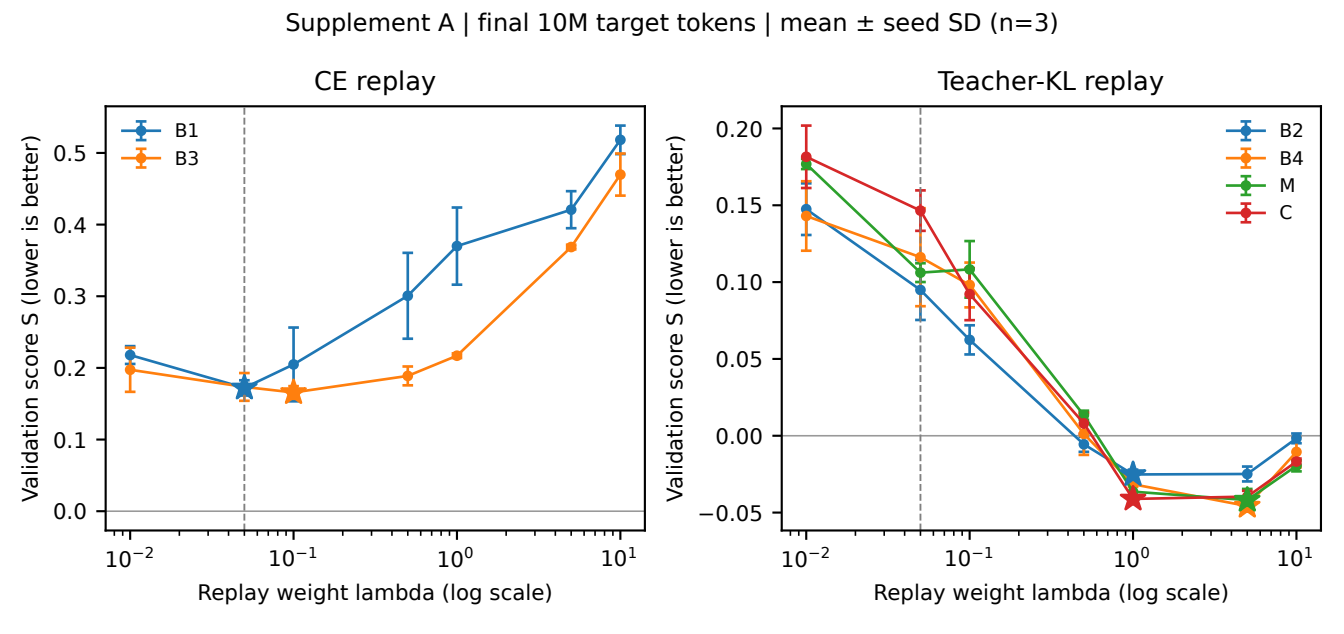
Supplement A | final 10M target tokens | mean ± seed SD (n=3)
CE replay
Teacher-KL replay
B1
B3
B2
B4
M
C
Validation score S (lower is better)
Replay weight lambda (log scale)
0.0
0.1
0.2
0.3
0.4
0.5
−0.05
0.00
0.05
0.10
0.15
0.20
10^{-2}
10^{-1}
10^{0}
10^{1}